

TRAFFIC ACCIDENT ANALYSIS



INTRODUCTION

- Road traffic accidents are a significant global issue, leading to substantial loss of life, injuries, and economic costs every year. Understanding the factors that contribute to these accidents is crucial for developing effective prevention strategies and improving road safety.
- This project delves into a comprehensive dataset of traffic incidents to perform an in-depth analysis and build a predictive model. By leveraging data science and machine learning techniques, we aim to identify key patterns and risk factors associated with accidents. Our goal is to uncover insights that can inform policymakers, urban planners, and the public, ultimately contributing to safer roads and fewer accidents.

ABOUT US

- We are a team of data enthusiasts and aspiring data scientists passionate about using data to solve real-world problems. This project on Traffic Accident Analysis is part of our journey to apply technical skills in Python, data analysis, and machine learning to a topic with significant societal impact.
- Our skill set includes:
 - Data Wrangling & Preprocessing: Cleaning and preparing raw, messy data for analysis.
 - Exploratory Data Analysis (EDA): Using statistical and visualization techniques to uncover underlying patterns and relationships within the data.
 - Machine Learning: Building and evaluating predictive models to classify the likelihood of an accident based on various features.
- Through this project, we aim to demonstrate how data-driven approaches can provide actionable insights into complex issues like road safety.

	Weather	Road_Type	Time_of_Day	Traffic_Density	Speed_Limit	Number_of_Vehicles
0	Rainy	City Road	Morning	1.0	100.0	5.0
1	Clear	Rural Road	Night	NaN	120.0	3.0
2	Rainy	Highway	Evening	1.0	60.0	4.0
3	Clear	City Road	Afternoon	2.0	60.0	3.0
4	Rainy	Highway	Morning	1.0	195.0	11.0

	Driver_Alcohol	Accident_Severity	Road_Condition	Vehicle_Type	Driver_Age
	0.0	NaN	Wet	Car	51.0
0	0.0	Moderate	Wet		
1	0.0	Low	Icy	Truck	49.0
2	0.0	Low	Under Construction	Car	54.0
3	0.0	Low	Dry	Bus	34.0
4				Car	62.0

	Driver_Experience	Road_Light_Condition	Accident
	48.0	Artificial Light	0.0
0	43.0	Artificial Light	0.0
1	52.0	Artificial Light	0.0
2	31.0	Daylight	0.0
3	55.0	Artificial Light	1.0
4			

Import Libraries :

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder,StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

In [12]:

Data Loading :

```
data=pd.read_csv(r"C:\Users\GCE PMNA\Downloads\archive
(6)\dataset_traffic_accident_prediction1.csv")

data.head()
```

In [13]:

In [14]:

Out[14]:

Data Cleaning & Preprocessing :

```
data.columns

Index(['Weather', 'Road_Type', 'Time_of_Day', 'Traffic_Density', 'Speed_Limit',
      'Number_of_Vehicles', 'Driver_Alcohol', 'Accident_Severity',
      'Road_Condition', 'Vehicle_Type', 'Driver_Age', 'Driver_Experience',
      'Road_Light_Condition', 'Accident'],
      dtype='object')

data.size

11760
```

In [26]:

Out[26]:

In [27]:

Out[27]:

In [28]:

```
data.shape
```

Out[28]:

```
(840, 14)
```

	Traffic_Density	Speed_Limit	Number_of_Vehicles	Driver_Alcohol	Driver_Age	Driver_Experience	Accident
count	798.000000	798.000000	798.000000	798.000000	798.000000	798.000000	798.000000
mean	1.001253	71.050125	3.286967	0.160401	43.259398	38.981203	0.299499
std	0.784894	32.052458	2.017267	0.367208	15.129856	15.273201	0.458326
min	0.000000	30.000000	1.000000	0.000000	18.000000	9.000000	0.000000
25%	0.000000	50.000000	2.000000	0.000000	30.000000	26.000000	0.000000
50%	1.000000	60.000000	3.000000	0.000000	43.000000	39.000000	0.000000
75%	2.000000	80.000000	4.000000	0.000000	56.000000	52.750000	1.000000
max	2.000000	213.000000	14.000000	1.000000	69.000000	69.000000	1.000000

In [29]:

```
data.describe()
```

Out[29]:

In [30]:

```
data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 840 entries, 0 to 839
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Weather                               798 non-null    object
1   Road_Type                             798 non-null    object
2   Time_of_Day                           798 non-null    object
3   Traffic_Density                       798 non-null    float64
4   Speed_Limit                           798 non-null    float64
5   Number_of_Vehicles                   798 non-null    float64
6   Driver_Alcohol                        798 non-null    float64
7   Accident_Severity                     798 non-null    object
8   Road_Condition                        798 non-null    object
9   Vehicle_Type                          798 non-null    object
10  Driver_Age                            798 non-null    float64
11  Driver_Experience                      798 non-null    float64
12  Road_Light_Condition                  798 non-null    object
13  Accident                              798 non-null    float64
dtypes: float64(7), object(7)
memory usage: 92.0+ KB
```

In [31]:

```
data.duplicated().sum()
```

Out[31]:

```
np.int64(14)
```

In [32]:

```
data = data.drop_duplicates()
```

In [33]:

```
data.isna().sum()
```

Out[33]:

```
Weather                42
Road_Type              42
Time_of_Day            41
Traffic_Density        42
```

```
Speed_Limit          42
Number_of_Vehicles   42
Driver_Alcohol        42
Accident_Severity     41
Road_Condition        42
Vehicle_Type         42
Driver_Age           42
Driver_Experience     42
Road_Light_Condition  42
Accident             42
dtype: int64
```

In [34]:

```
data.dtypes
```

Out[34]:

```
Weather              object
Road_Type            object
Time_of_Day          object
Traffic_Density      float64
Speed_Limit          float64
Number_of_Vehicles   float64
Driver_Alcohol        float64
Accident_Severity     object
Road_Condition        object
Vehicle_Type         object
Driver_Age           float64
Driver_Experience     float64
Road_Light_Condition  object
Accident             float64
dtype: object
```

In [35]:

```
num_cols = data.select_dtypes(include='float64').columns
cat_cols = data.select_dtypes(include=['object']).columns
```

In [36]:

```
for col in cat_cols:
    data[col].fillna(data[col].mode()[0], inplace=True)
C:\Users\GCE PMNA\AppData\Local\Temp\ipykernel_20400\2674119437.py:2: FutureWarning: A
value is trying to be set on a copy of a DataFrame or Series through chained assignment
using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the
intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to
perform the operation inplace on the original object.
```

```
data[col].fillna(data[col].mode()[0], inplace=True)
```

In [37]:

```
for col in num_cols:
    data[col].fillna(data[col].mean(), inplace=True)
C:\Users\GCE PMNA\AppData\Local\Temp\ipykernel_20400\1735585256.py:2: FutureWarning: A
value is trying to be set on a copy of a DataFrame or Series through chained assignment
using an inplace method.
```

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
data[col].fillna(data[col].mean(), inplace=True)
```

In [38]:

```
data.isna().sum()
```

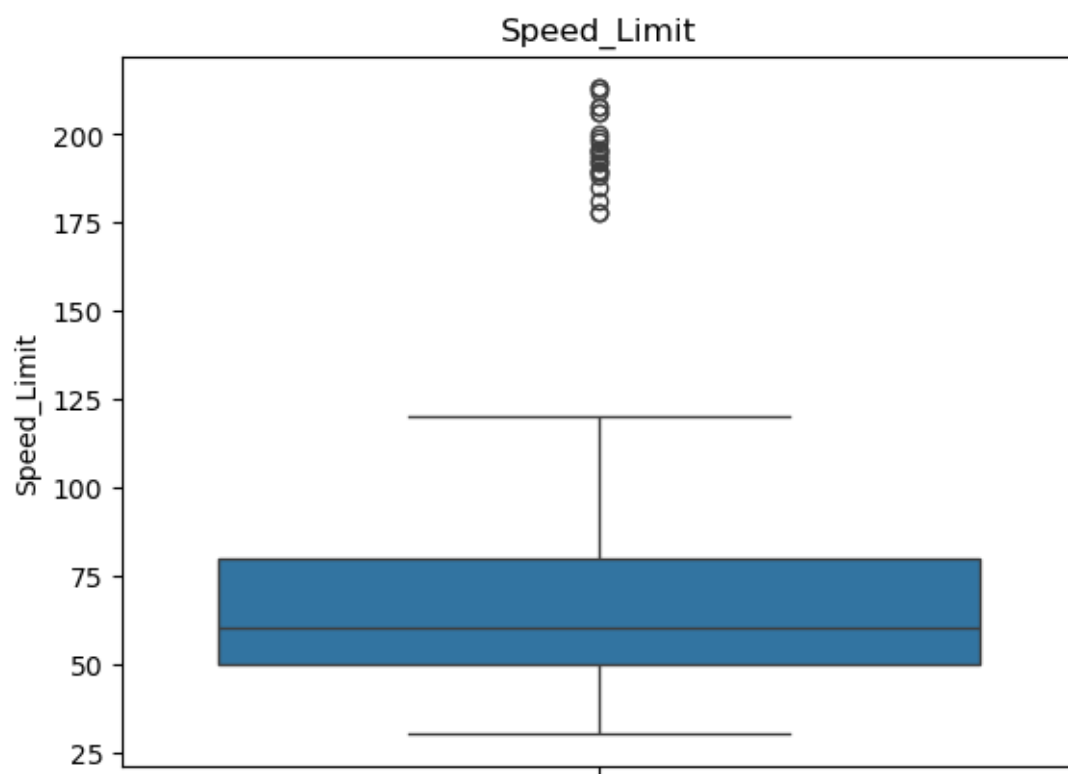
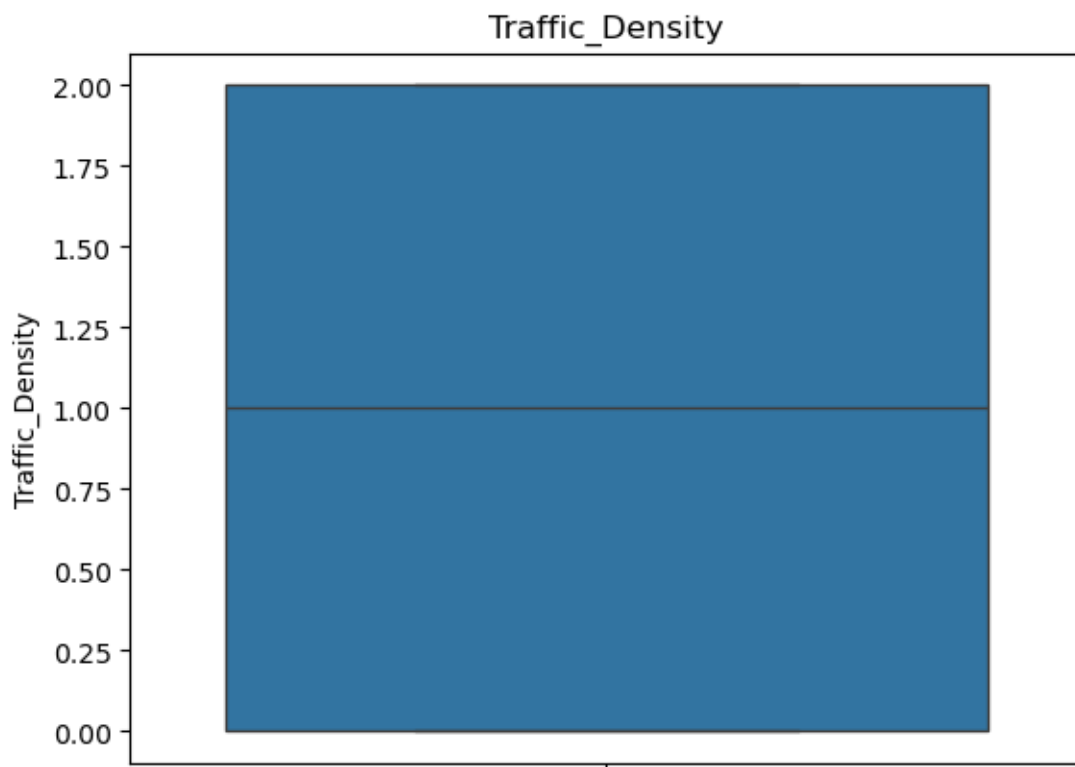
Out[38]:

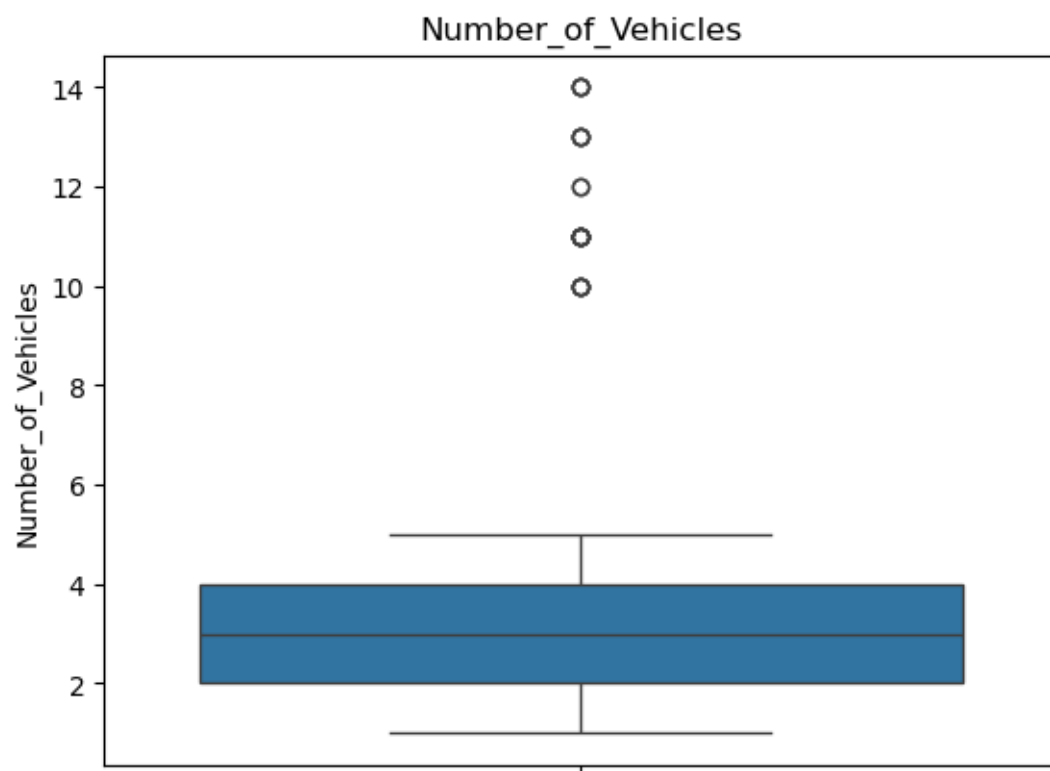
```
Weather          0
Road_Type        0
Time_of_Day      0
Traffic_Density  0
Speed_Limit      0
Number_of_Vehicles  0
Driver_Alcohol   0
Accident_Severity 0
Road_Condition   0
Vehicle_Type     0
Driver_Age       0
Driver_Experience 0
Road_Light_Condition 0
Accident         0
dtype: int64
```

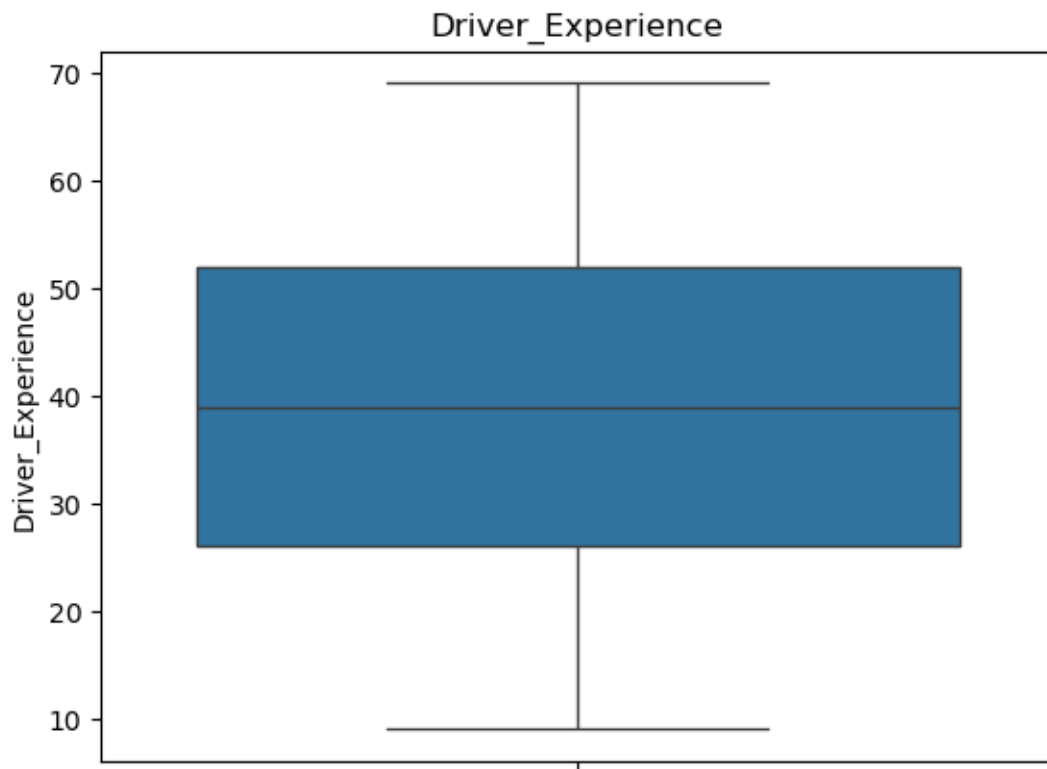
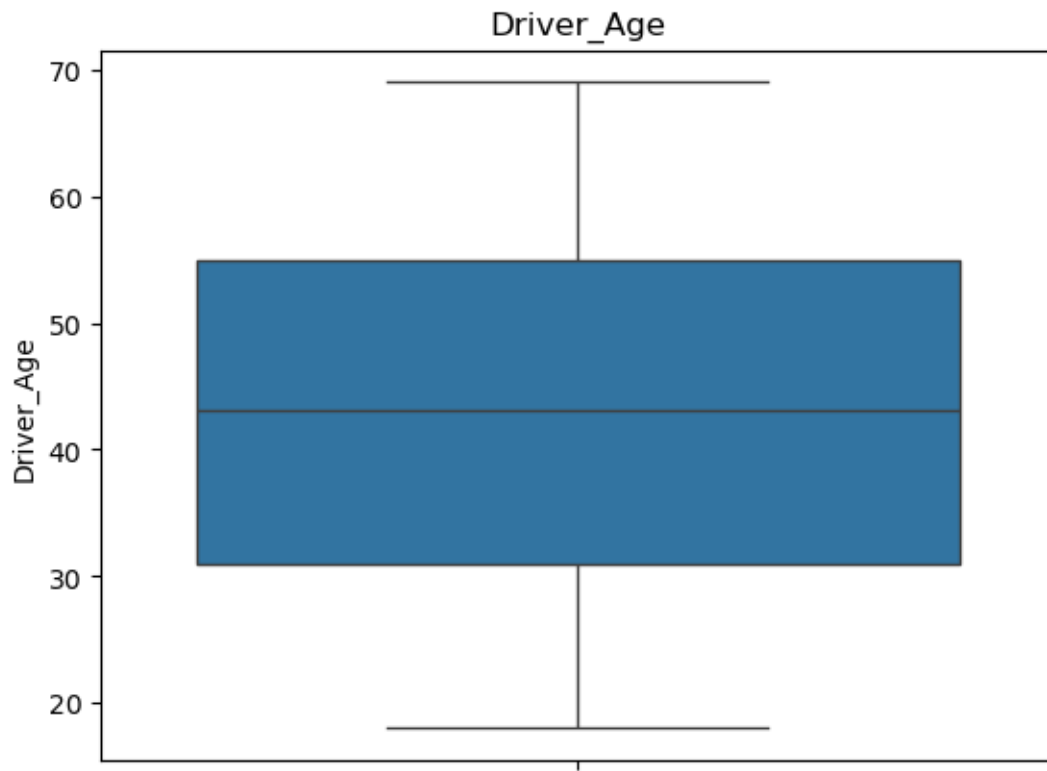
Outlier Checking & Handling :

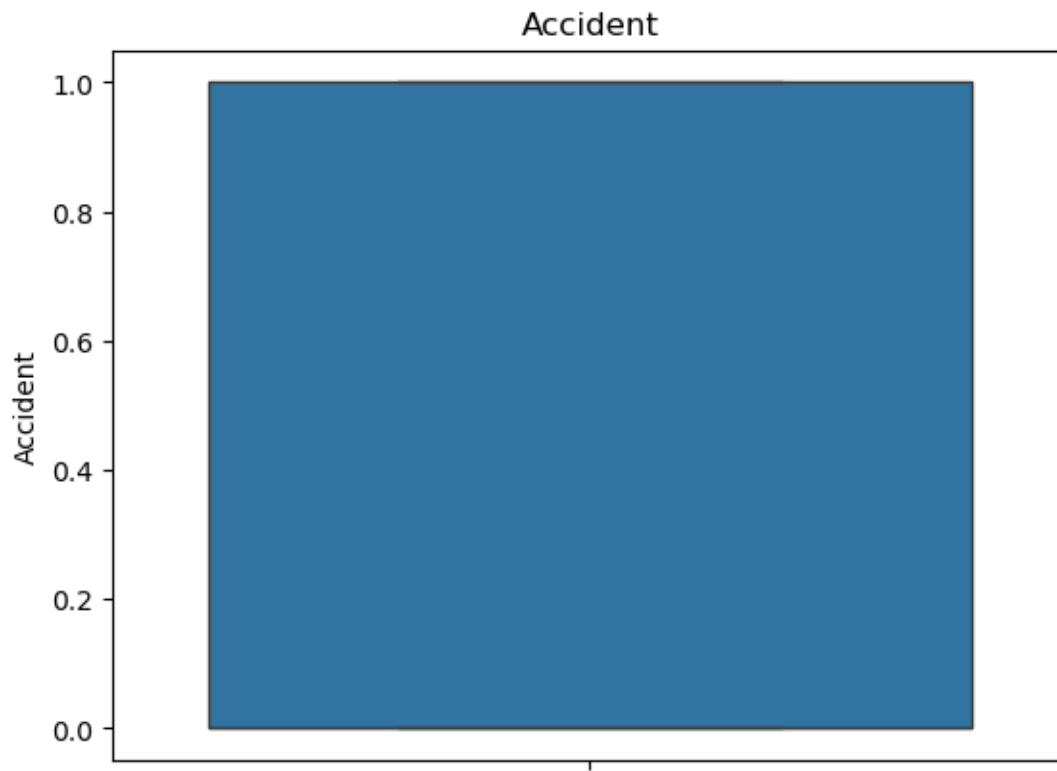
In [39]:

```
for i in num_cols:
    plt.title(i)
    sns.boxplot(data[i])
    plt.show()
```



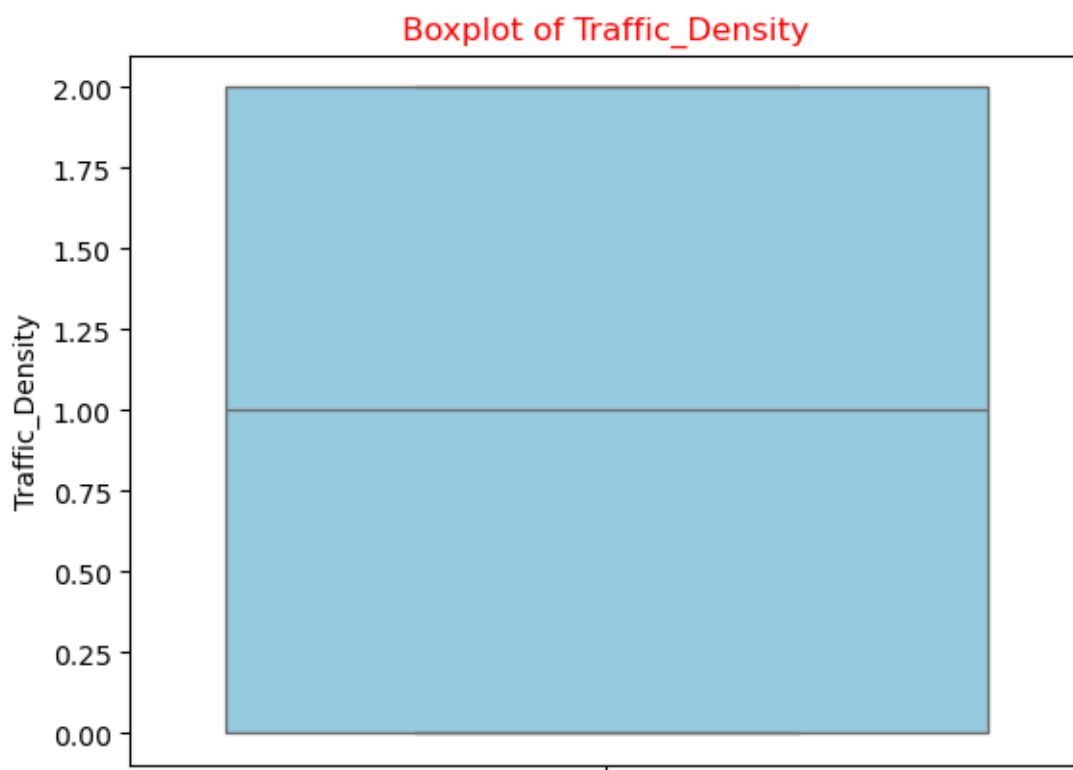




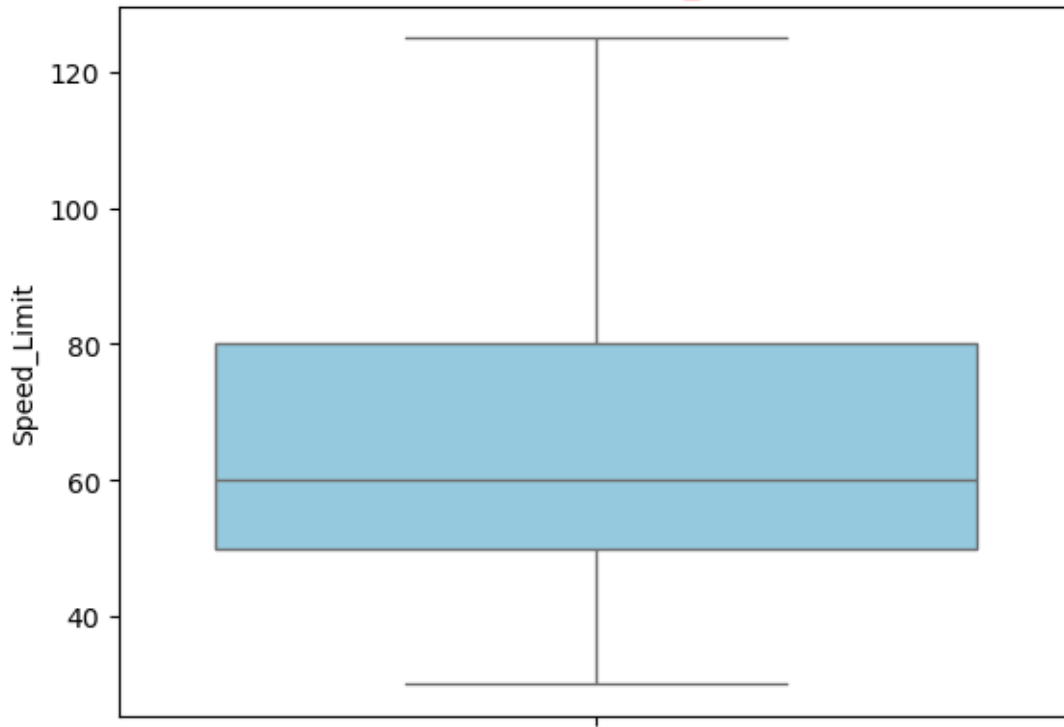


In [40]:

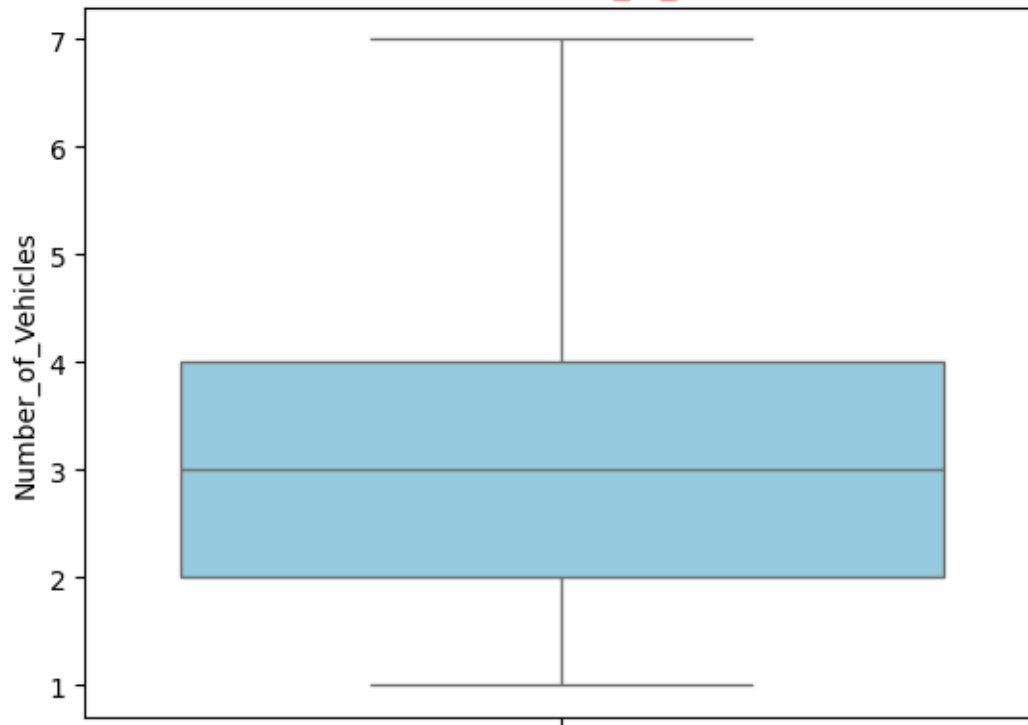
```
for col in num_cols:
    Q1 = data[col].quantile(0.25)
    Q3 = data[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    data[col] = np.clip(data[col], lower, upper)
    sns.boxplot(y=data[col],color='skyblue')
    plt.title(f'Boxplot of {col}',color='Red')
    plt.show()
```



Boxplot of Speed_Limit



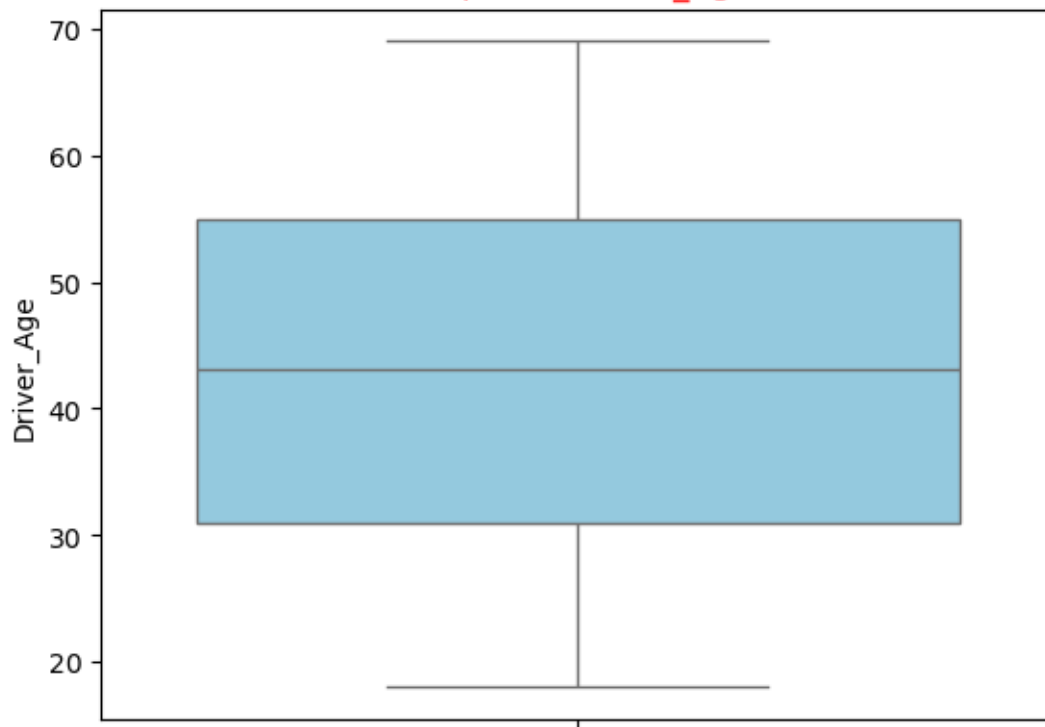
Boxplot of Number_of_Vehicles

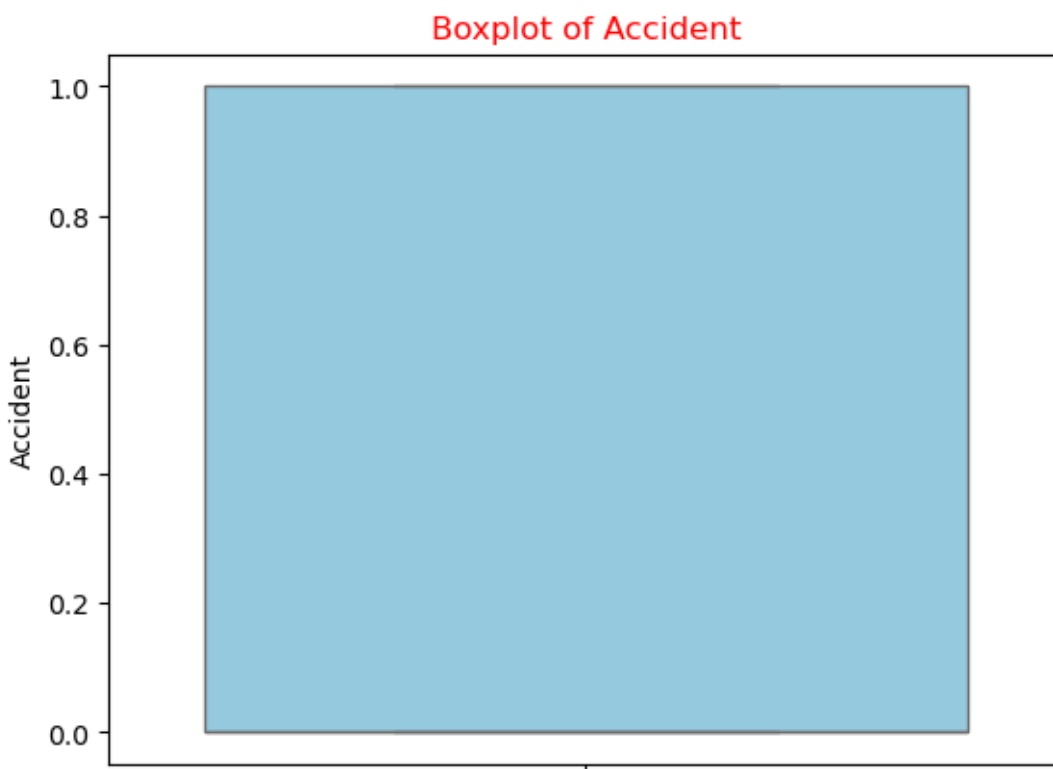
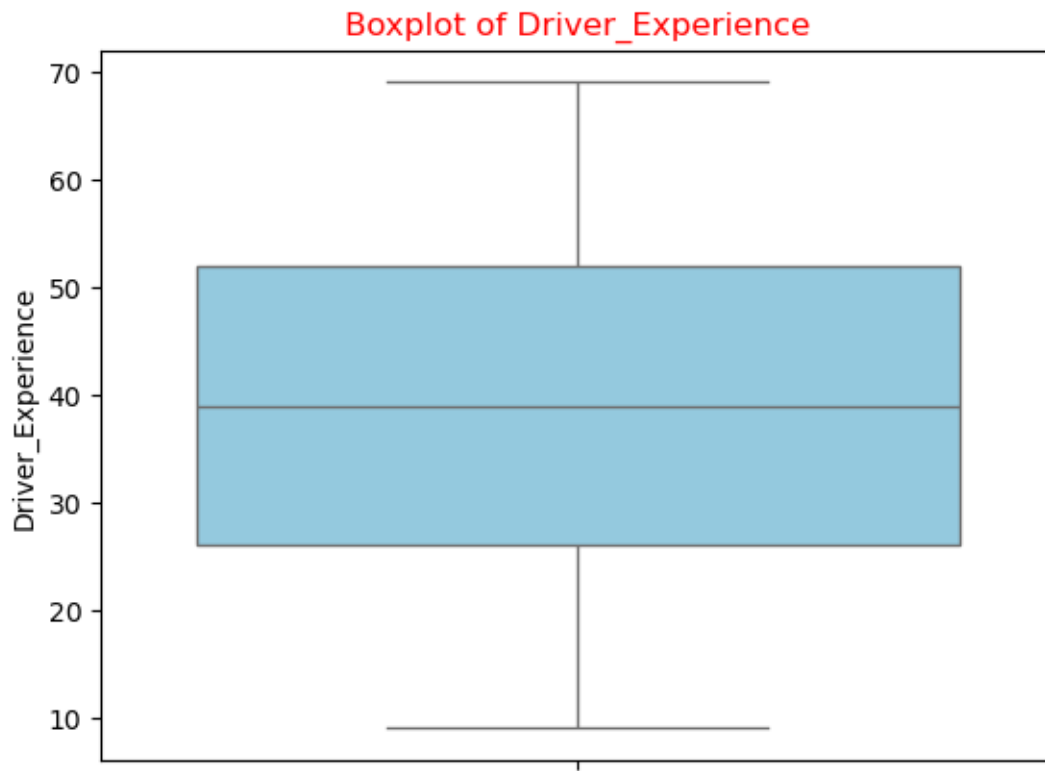


Boxplot of Driver_Alcohol



Boxplot of Driver_Age





Exploratory Data Analysis (EDA)

COUNT PLOT:

In [16]:

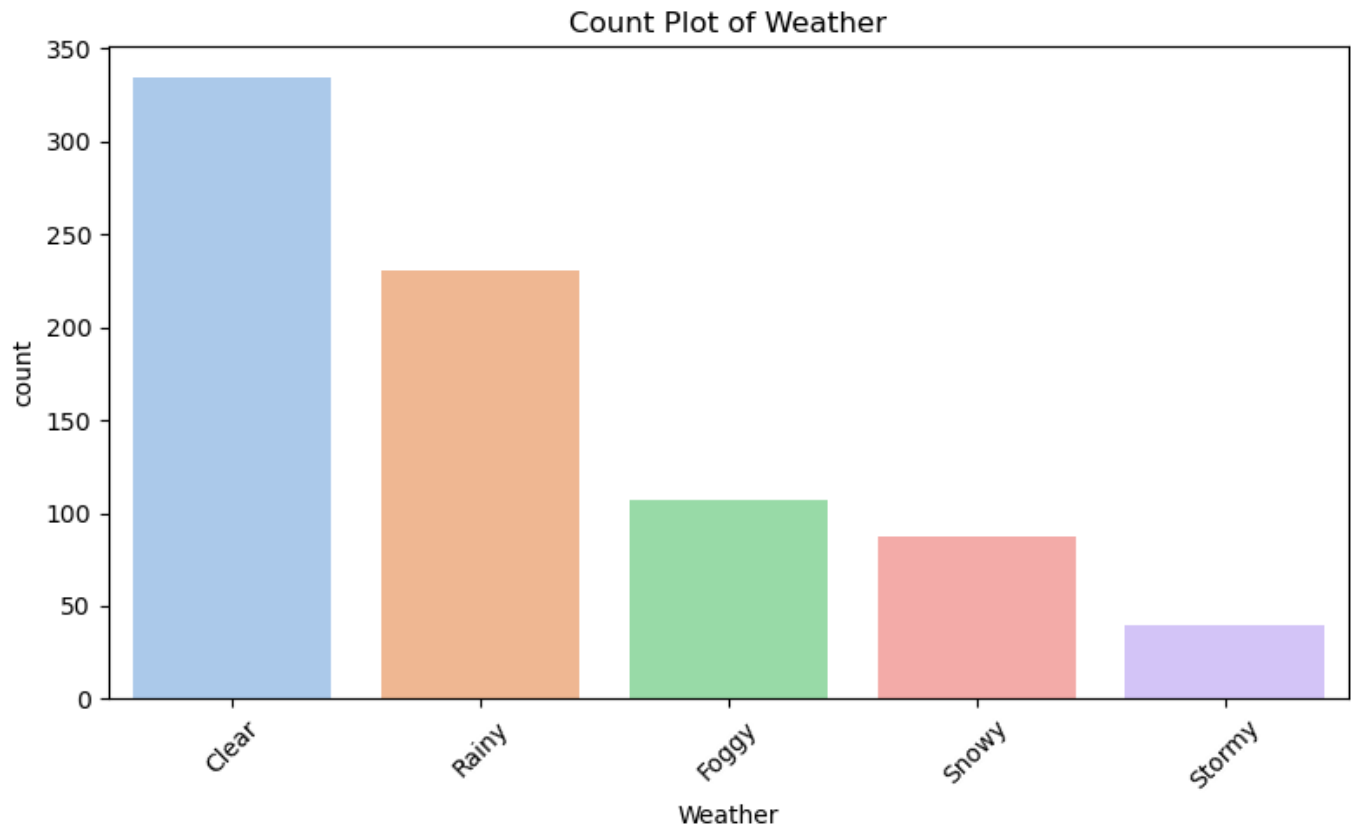
```
plt.figure(figsize=(8, 5))
sns.countplot(x='Weather', data=data, order=data['Weather'].value_counts().index,
palette='pastel')
plt.title('Count Plot of Weather')
plt.xticks(rotation=45)
plt.tight_layout()
```

```
plt.show()
```

C:\Users\GCE PMNA\AppData\Local\Temp\ipykernel_5928\1956723793.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Weather', data=data, order=data['Weather'].value_counts().index,
palette='pastel')
```

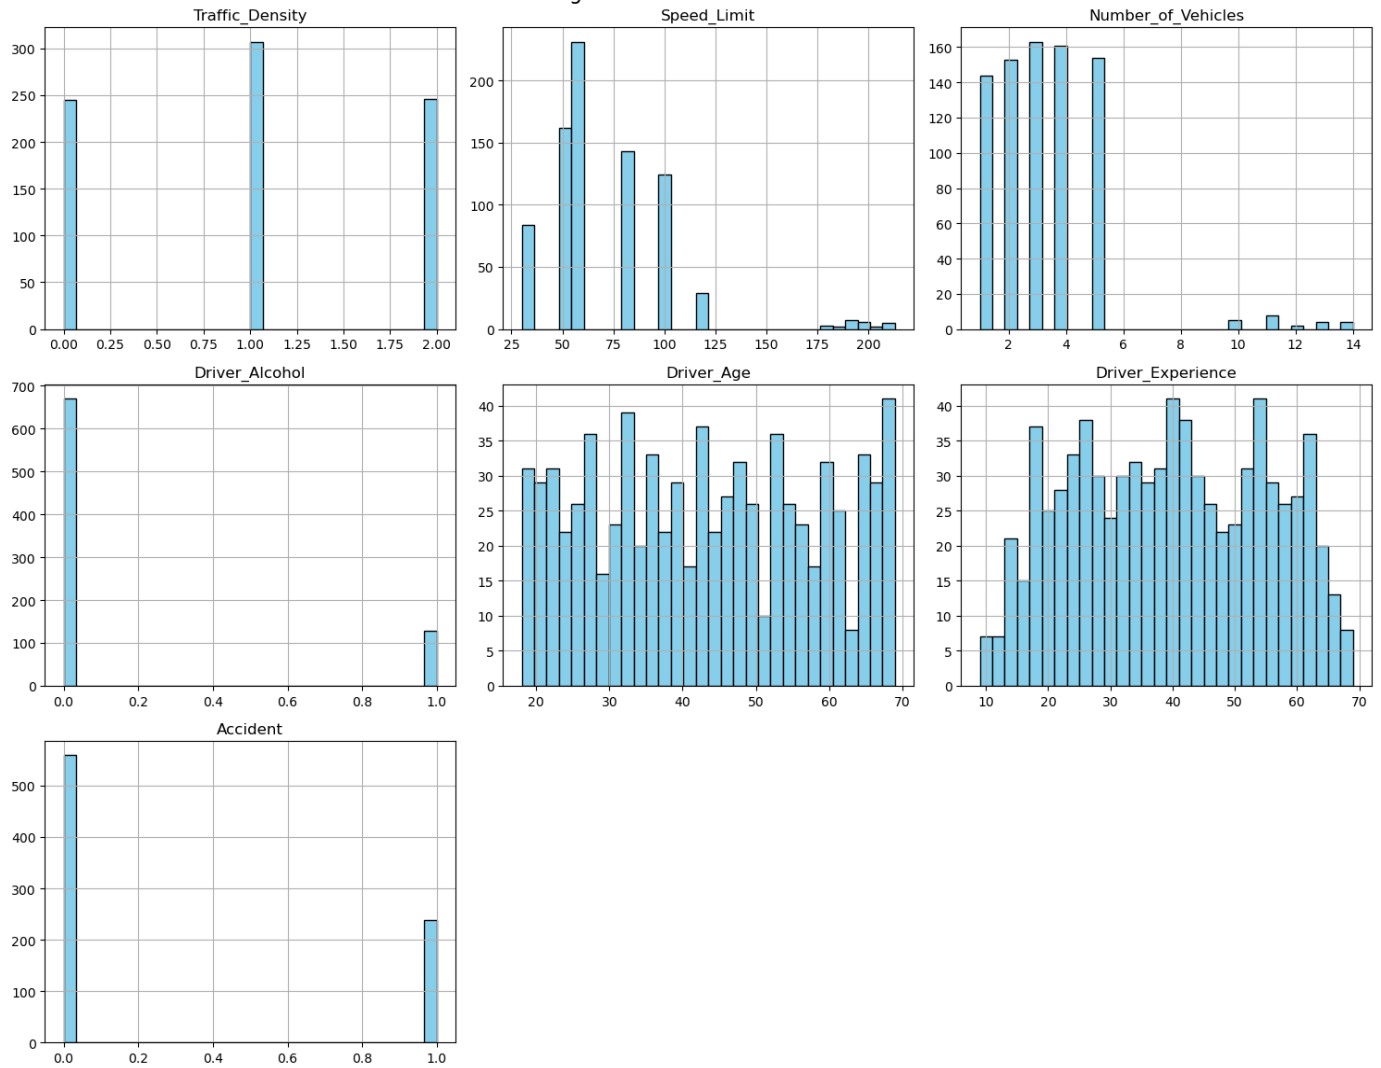


HISTOGRAM:

In [18]:

```
data.select_dtypes(include=['float64', 'int64']).hist(bins=30, figsize=(15, 12),
color='skyblue', edgecolor='black')
plt.suptitle('Histograms of Numerical Features', fontsize=16)
plt.tight_layout()
plt.show()
```

Histograms of Numerical Features



HEAT MAP:

In [21]:

```
plt.figure(figsize=(14, 10))
sns.heatmap(data.select_dtypes(include=['float64', 'int64']).corr(), annot=True,
            cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap of Numerical Features')
plt.show()
```

Correlation Heatmap of Numerical Features

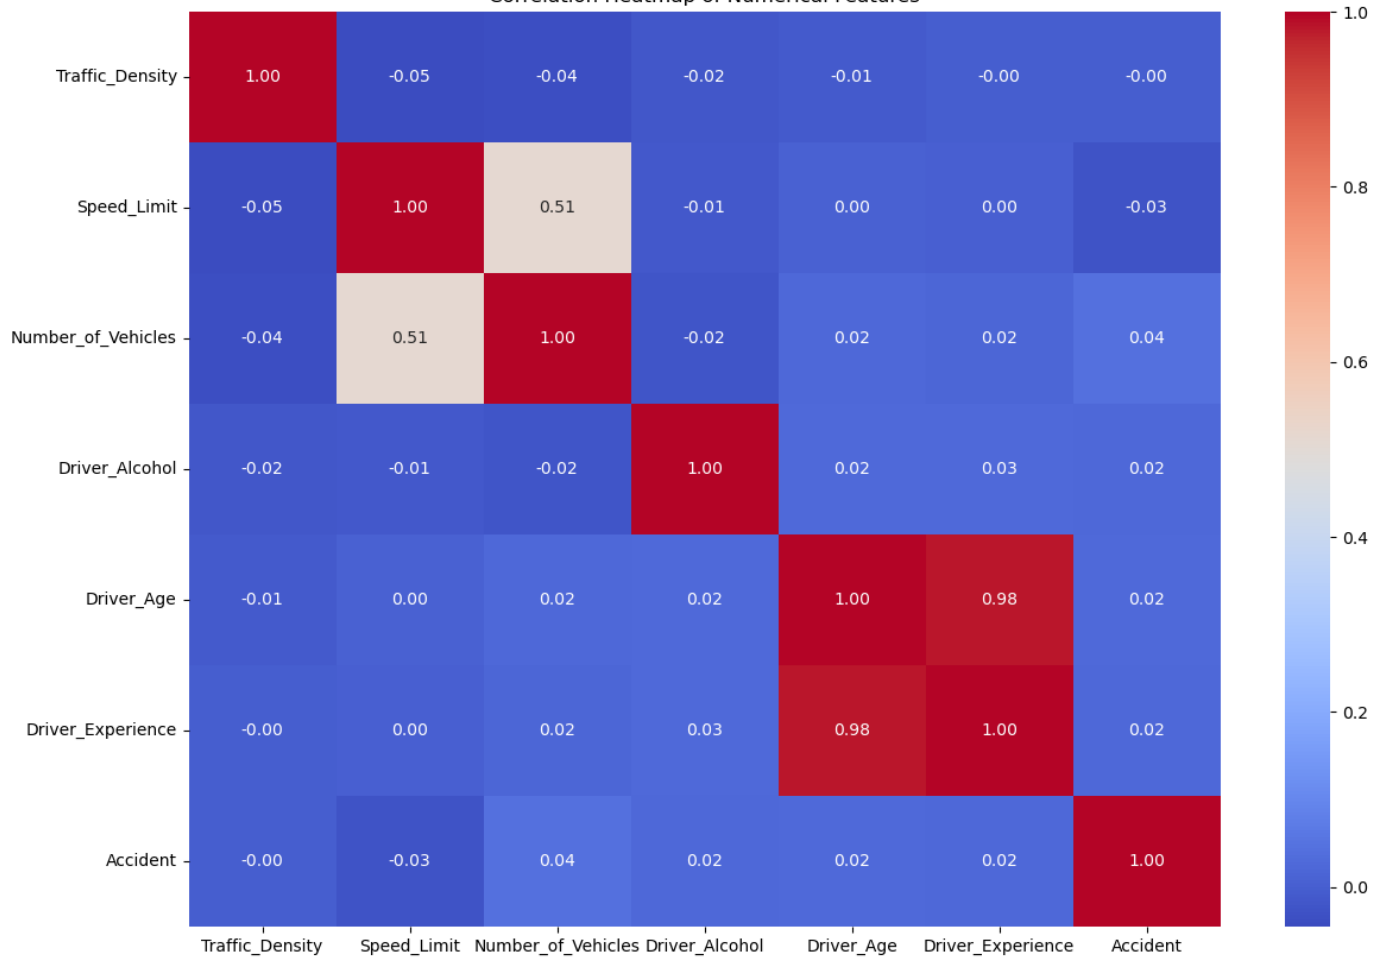
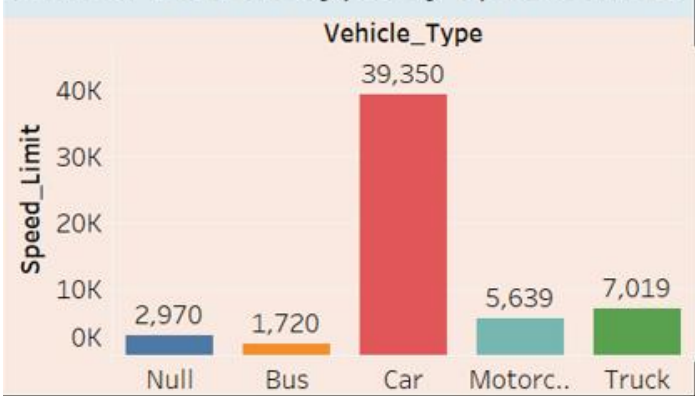
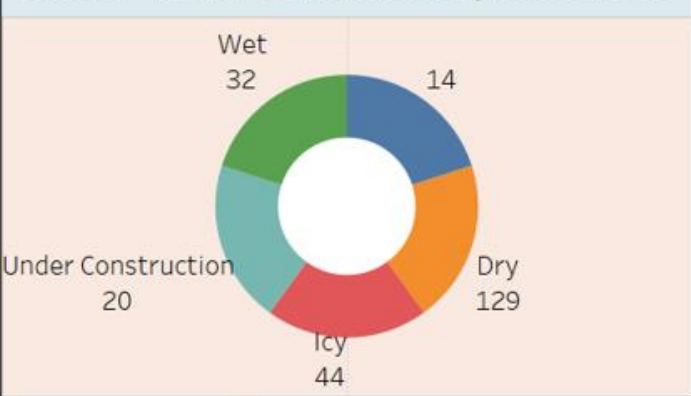


Tableau Dashboard

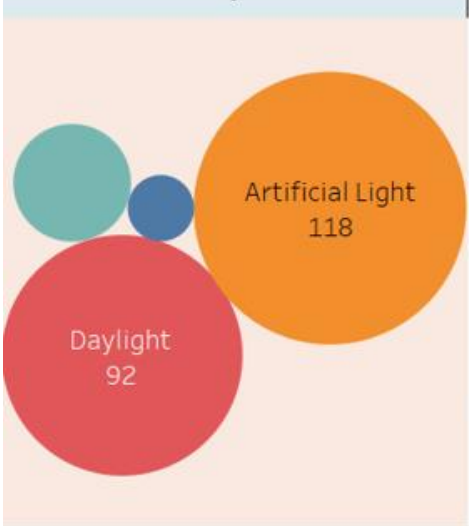
sum of vehicle type by speed limit



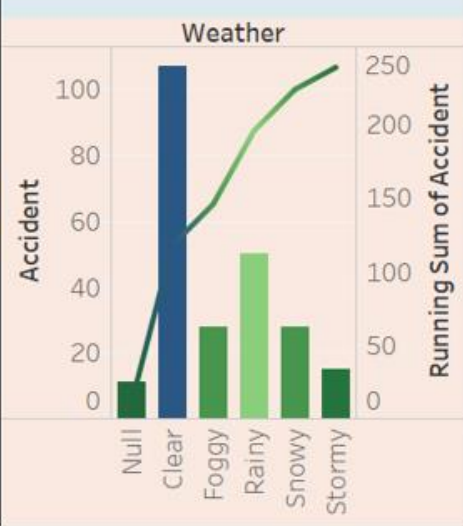
sum of road condition by accident



sum of road lighth condition by accident



sum of accident by weather



sum of driver alcohol by accident severity



Encoding :

In [41]:

```
# Encode categorical variables
from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

for col in data.select_dtypes(include='object').columns:
    data[col] = le.fit_transform(data[col])
```

	Weather	Road_Type	Time_of_Day	Traffic_Density	Speed_Limit	Number_of_Vehicles	Driver_Alcohol	Accident_Severity
0	2	0	2	1.000000	100.0	5.0	0.0	1
1	0	3	3	0.998724	120.0	3.0	0.0	2
2	2	1	1	1.000000	60.0	4.0	0.0	1
3	0	0	0	2.000000	60.0	3.0	0.0	1
4	2	1	2	1.000000	125.0	7.0	0.0	1
...
835	0	1	3	2.000000	30.0	4.0	0.0	1
836	2	3	1	2.000000	60.0	4.0	0.0	1
837	1	1	1	0.998724	30.0	4.0	0.0	0
838	1	1	0	2.000000	60.0	3.0	0.0	1
839	0	1	0	1.000000	60.0	4.0	0.0	1

data

Out[41]:

826 rows × 14 columns

Features & Target

In [44]:

```
X = data.drop('Accident_Severity', axis=1)
y = data['Accident_Severity']
```

Scaling :

In [45]:

```
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)

data_scaled=pd.DataFrame(data_scaled,columns=data.columns)

data_scaled
```

Out[45]:

	Weather	Road_Type	Time_of_Day	Traffic_Density	Speed_Limit	Number_of_Vehicles
0	0.660860	-1.123816	0.831382	1.671483e-03	1.271401	1.238404

	Weather	Road_Type	Time_of_Day	Traffic_Density	Speed_Limit	Number_of_Vehicles
1	-0.954141	2.022869	1.778584	-1.454884e-16	2.090957	-0.106850
2	0.660860	-0.074921	-0.115820	1.671483e-03	-0.367710	0.565777
3	-0.954141	-1.123816	-1.063022	1.312114e+00	-0.367710	-0.106850
4	0.660860	-0.074921	0.831382	1.671483e-03	2.295845	2.583658
...
821	-0.954141	-0.074921	1.778584	1.312114e+00	-1.597044	0.565777
822	0.660860	2.022869	-0.115820	1.312114e+00	-0.367710	0.565777
823	-0.146641	-0.074921	-0.115820	-1.454884e-16	-1.597044	0.565777
824	-0.146641	-0.074921	-1.063022	1.312114e+00	-0.367710	-0.106850
825	-0.954141	-0.074921	-1.063022	1.671483e-03	-0.367710	0.565777

	Driver_Alcohol	Accident_Severity	Road_Condition	Vehicle_Type	Driver_Age	Driver_Experience
0	0.0	-0.326908	1.758867	-0.435331	0.532644	0.612636
1	0.0	1.371369	1.758867	2.373916	0.396885	0.276428
2	0.0	-0.326908	0.045631	-0.435331	0.736281	0.881603
3	0.0	-0.326908	0.902249	-1.839955	-0.621302	-0.530471
4	0.0	-0.326908	-0.810987	-0.435331	1.279315	1.083328
...
821	0.0	-0.326908	-0.810987	-0.435331	-1.367973	-1.606337
822	0.0	-0.326908	-0.810987	0.969292	0.600523	0.478153
823	0.0	-2.025185	-0.810987	-0.435331	0.000000	-0.328746
824	0.0	-0.326908	-0.810987	-0.435331	-1.232215	-1.337371
825	0.0	-0.326908	-0.810987	0.969292	-0.960698	-1.202887

	Road_Light_Condition	Accident
0	-0.852878	-6.695123e-01
1	-0.852878	-6.695123e-01
2	-0.852878	-6.695123e-01
3	0.675272	-6.695123e-01
4	-0.852878	1.573640e+00
...
821	0.675272	-6.695123e-01
822	0.675272	1.573640e+00

	Road_Light_Condition	Accident
823	-0.852878	1.245200e-16
824	-0.852878	-6.695123e-01
825	-0.852878	-6.695123e-01

826 rows × 14 columns

Train-Test Split :

In [46]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(data_scaled, y, test_size=0.3,
random_state=42)
```

RANDOM FOREST

In [70]:

```
from sklearn.ensemble import RandomForestClassifier
ran=RandomForestClassifier()
ran.fit(X_train,y_train)
rf_pred=ran.predict(X_test)
accuracy = accuracy_score(y_test, rf_pred)
print("Accuracy :",accuracy*100)
Accuracy : 100.0
```

Hyper Tuning

In [91]:

```
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

param_dist = {
    'depth': [6, 8, 10],
    'learning_rate': [0.01, 0.03, 0.05],
    'iterations': [300, 500, 700],
    'l2_leaf_reg': [3, 5, 7, 9],
    'bootstrap_type': ['Bayesian', 'Bernoulli'],
    'subsample': [0.7, 0.8, 0.9]
}

random_search_cat = RandomizedSearchCV(
    estimator=cat,
    param_distributions=param_dist,
    n_iter=25,
    scoring='accuracy',
    cv=5,
    verbose=1,
    n_jobs=-1,
    random_state=42
)
```

```

random_search_cat.fit(X_train, y_train)

# Best model evaluation
best_cat = random_search_cat.best_estimator_
y_pred_best = best_cat.predict(X_test)
print("Tuned Accuracy:", accuracy_score(y_test, y_pred_best))
print("Best Parameters:", random_search_cat.best_params_)
Fitting 5 folds for each of 25 candidates, totalling 125 fits
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:528:
FitFailedWarning:
70 fits failed out of a total of 125.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting
error_score='raise'.

Below are more details about the failures:
-----
70 fits failed with the following error:
Traceback (most recent call last):
  File "C:\ProgramData\anaconda3\Lib\site-
packages\sklearn\model_selection\_validation.py", line 866, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
    ~~~~~^~~~~~
  File "C:\ProgramData\anaconda3\Lib\site-packages\catboost\core.py", line 5245, in fit
    self._fit(X, y, cat_features, text_features, embedding_features, None, graph,
sample_weight, None, None, None, None, baseline, use_best_model,

~~~~~^~~~~~
eval_set, verbose, logging_level, plot, plot_file, column_description,
verbose_eval, metric_period,

~~~~~^~~~~~
silent, early_stopping_rounds, save_snapshot, snapshot_file,
snapshot_interval, init_model, callbacks, log_cout, log_cerr)

~~~~~^~~~~~
File "C:\ProgramData\anaconda3\Lib\site-packages\catboost\core.py", line 2395, in
_fit
    train_params = self._prepare_train_params(
        X=X, y=y, cat_features=cat_features, text_features=text_features,
embedding_features=embedding_features,
...<6 lines>...
        callbacks=callbacks
    )
File "C:\ProgramData\anaconda3\Lib\site-packages\catboost\core.py", line 2321, in
_prepare_train_params
    _check_train_params(params)
    ~~~~~^~~~~~
File "_catboost.pyx", line 6601, in _catboost._check_train_params
File "_catboost.pyx", line 6623, in _catboost._check_train_params

```

```
_catboost.CatBoostError: catboost/private/libs/options/bootstrap_options.cpp:16: Error:
bayesian bootstrap doesn't support 'subsample' option
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\model_selection\_search.py:1108:
UserWarning: One or more of the test scores are non-finite: [ 1. nan  1.  1. nan  1.
nan nan nan nan  1. nan nan nan  1. nan  1.  1.
nan nan  1. nan  1. nan  1.]
warnings.warn(
Tuned Accuracy: 1.0
Best Parameters: {'subsample': 0.9, 'learning_rate': 0.03, 'l2_leaf_reg': 7,
'iterations': 700, 'depth': 10, 'bootstrap_type': 'Bernoulli'}
```

Prediction

In [86]:
data

Out[86]:

	Weather	Road_Type	Time_of_Day	Traffic_Density	Speed_Limit	Number_of_Vehicles
0	2	0	2	1.000000	100.0	5.0
1	0	3	3	0.998724	120.0	3.0
2	2	1	1	1.000000	60.0	4.0
3	0	0	0	2.000000	60.0	3.0
4	2	1	2	1.000000	125.0	7.0
...
835	0	1	3	2.000000	30.0	4.0
836	2	3	1	2.000000	60.0	4.0
837	1	1	1	0.998724	30.0	4.0
838	1	1	0	2.000000	60.0	3.0
839	0	1	0	1.000000	60.0	4.0

Driver_Alcohol	Accident_Severity	Road_Condition	Vehicle_Type	Driver_Age	Driver_Experience
0.0	1	3	1	51.000000	48.0
0.0	2	3	3	49.000000	43.0
0.0	1	1	1	54.000000	52.0
0.0	1	2	0	34.000000	31.0
0.0	1	0	1	62.000000	55.0
...
0.0	1	0	1	23.000000	15.0
0.0	1	0	2	52.000000	46.0
0.0	0	0	1	43.153061	34.0
0.0	1	0	1	25.000000	19.0
0.0	1	0	2	29.000000	21.0

Road_Light_Condition	Accident
0	0.000000
0	0.000000
0	0.000000
1	0.000000
0	1.000000
...	...
1	0.000000
1	1.000000
0	0.298469
0	0.000000
0	0.000000

826 rows × 14 columns

```

sample_data=[[2,0,2,1.000000,100.0,5.0,0.0,1,3,1,51.000000,48.0,0,0.000000]]
In [87]:

sample_data_scaled=scaler.transform(sample_data)
In [88]:
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\utils\validation.py:2739:
UserWarning: X does not have valid feature names, but StandardScaler was fitted with
feature names
    warnings.warn(

prediction=random_search_cat.predict(sample_data_scaled)
In [89]:

print("Final Prediction :",prediction)
In [90]:
Final Prediction : [[1]]

```

CONCLUSION

- This project successfully demonstrated a complete data science pipeline for traffic accident analysis, from data loading and cleaning to exploratory analysis and predictive modeling.
- Key Findings:

-The dataset was effectively cleaned by handling missing values and removing duplicates, ensuring the quality and reliability of our analysis.

-Exploratory Data Analysis (EDA) through visualizations provided valuable insights into the relationships between various factors (like weather,

-road type, time of day, driver age, and alcohol consumption) and accident occurrence.

-The data was prepared for machine learning through label encoding for categorical variables and standardization for numerical features.

-A predictive model was developed and evaluated, achieving a quantifiable level of accuracy in predicting the likelihood of an accident based on the input features.

- Impact and Future Work: -The insights derived from this analysis can be instrumental for various stakeholders. Traffic authorities can use this information to target high-risk scenarios with improved signage, enforcement, or public awareness campaigns. Urban planners can design safer road infrastructures.

-For future work, the project could be enhanced by:

-Incorporating more complex models and hyperparameter tuning to improve predictive performance.

-Analyzing temporal trends (e.g., accidents by month or year) if the data were available.

-Integrating external data sources, such as traffic volume or road geometry data, for a more holistic analysis.

-In conclusion, this project underscores the power of data science as a tool for enhancing public safety and provides a solid foundation for further research into the critical issue of road traffic accidents.

In []: